

# Introduction to RL – Part 2

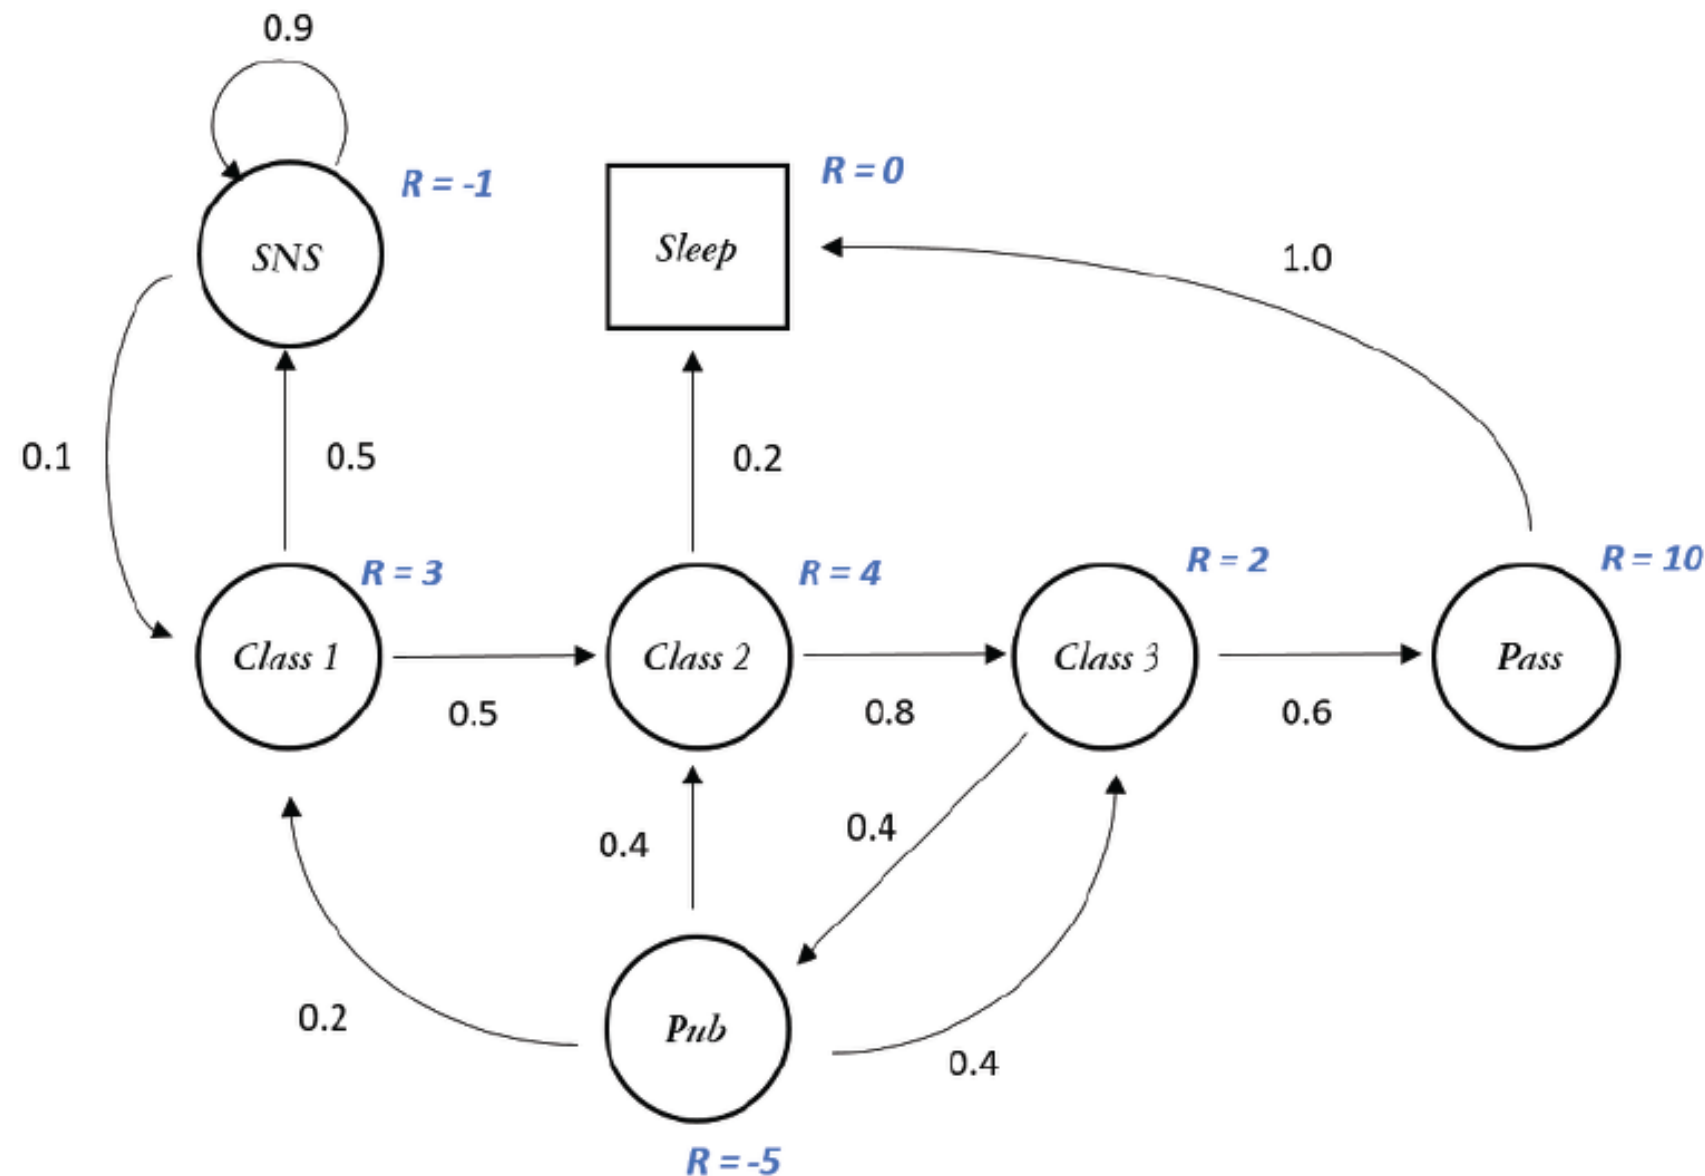
---

TA. Bogyeeong Suh

## Recap of Part 1

### Markov Decision Process

- The process of RL is usually given as **Markov Decision Processes (MDPs)**, which is a mathematical framework used for modeling decision making in situations where the outcomes are partly random and partly under the control of a decision maker.





Markov Reward Process (MRP)

# Recap of Part 1

## Policy Iteration

### 1) Policy Evaluation

-Given an arbitrary policy, we calculate value function  $V$  for **all states** under this policy. We iterate through each state and update  $V$  until it converges. → We find true value function with iteration



# Recap of Part 1

## Value Iteration

k = 0

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

0	0	0	0
0	0	0	0
0	0	0	0
0	0	0	0

at (1,2) state : state 1

Up  $V_1(s) = -1 + 0$   
 Down  $V_1(s) = -1 + 0$   
 Left  $V_1(s) = -1 + 0$   
 Right  $V_1(s) = -1 + 0$

$\therefore V_1(1) = \max V_1(s) = -1$

k = 1

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	0

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	0

at (1,2) state : state 1

Up  $V_2(s) = -1 + (-1)$   
 Down  $V_2(s) = -1 + (-1)$   
 Left  $V_2(s) = -1 + (0)$   
 Right  $V_2(s) = -1 + (-1)$

$\therefore V_2(s) = \max V_2(s) = -1$

0	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	-1
-1	-1	-1	0

at (1,3) state : state 2

Up  $V_2(s) = -1 + (-1)$   
 Down  $V_2(s) = -1 + (-1)$   
 Left  $V_2(s) = -1 + (-1)$   
 Right  $V_2(s) = -1 + (-1)$

$\therefore V_2(s) = \max V_2(s) = -2$

0	-1	-2	-2
-1	-2	-2	-2
-2	-2	-2	-1
-2	-2	-1	0

k = 2

0	-1	-2	-3
-1	-2	-3	-2
-2	-3	-2	-1
-3	-2	-1	0

k = 3

0	-1	-2	-3
-1	-2	-3	-2
-2	-3	-2	-1
-3	-2	-1	0


k =  $\infty$



# Monte Carlo method

## Monte Carlo method


- Generally, MC is a method that relies on **repeated random sampling** to obtain numerical results
- MC methods learn directly from episodes of **experience**
- In MC, we learn from *complete* episodes: no bootstrapping
- MC is *model-free*: no knowledge of MDP transitions / rewards




e.g. Finding the area of a circle with radius=1 using random samples of (x,y)

# Monte Carlo method

## Monte Carlo method



$$\begin{array}{ll}
 G_0 = R_1 + \gamma G_1 & G_5 = 0 \\
 G_1 = R_2 + \gamma G_2 & G_4 = R_5 + \gamma G_5 = 2 + 0.5 * 0 = 2 \\
 G_2 = R_3 + \gamma G_3 & G_3 = R_4 + \gamma G_4 = 1 + 0.5 * 2 = 2 \\
 G_3 = R_4 + \gamma G_4 & \rightarrow G_2 = R_3 + \gamma G_3 = 7 + 0.5 * 2 = 8 \\
 G_4 = R_5 + \gamma G_5 & G_1 = R_2 + \gamma G_2 = 4 + 0.5 * 8 = 8 \\
 G_5 = 0 & G_0 = R_1 + \gamma G_1 = 3 + 0.5 * 8 = 7
 \end{array}$$



$$\begin{aligned}
 G(s_1) &= R_1 + \gamma R_2 + \gamma^2 R_3 + \dots \\
 G(s_2) &= R_2 + \gamma^2 R_3 + \dots \\
 G(s_3) &= R_3 + \dots
 \end{aligned}$$

<Fig 2. The returns of each states in 1 episode>

- Goal: learn a policy that maximizes the total cumulative reward it receives over time
- Transition probability:  $p(s', r | s, \pi(s))$
- State Value function:  $V_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s]$  → **Average values of returns from episodes**
- Action value function:  $q_\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$

# Monte Carlo method

## Monte Carlo method

### 1. Initialization

**Initialize:**

$\pi(s) \in A(s), \forall s \in \mathbb{S}$

$Q(s, a) \in \mathbb{R}, \forall s \in \mathbb{S}, \forall a \in A(s)$

$Returns(s, a) \leftarrow \text{empty list}, \forall s \in \mathbb{S}, \forall a \in A(s)$

### 2. Monte Carlo

**Loop forever**(for each episode):

Choose  $S_0 \in \mathbb{S}, A_0 \in A(S_0)$ , randomly such that all pairs have probability  $> 0$   
 Generate an episode from  $S_0, A_0$  following  $\pi: S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$

$G \leftarrow 0$

**Loop** for each step of episode,  $t = T - 1, T - 2, \dots, 0$

$G \leftarrow \gamma G + R_{t+1}$

Unless the  $S_t, A_t$  appears in  $S_0, A_0, S_1, A_1, \dots, S_{t-1}, A_{t-1}$ :

Append  $G$  to  $Returns(S_t, A_t)$

$Q(S_t, A_t) \leftarrow \text{Average}(Returns(S_t, A_t))$

$\pi(S_t) \leftarrow \text{argmax}_a Q(S_t, a)$

- Using randomly generated episodes, calculate action value function from cumulative reward  $G$

$$q_{\pi}(s, a) = \frac{1}{N(s)} \sum_{i=1}^{N(s)} G_i(s)$$

- $N(s)$ : number of times we visited state  $s$  we from start to end in total episodes
- $G_i(s)$ : return of state  $s$  in episode  $i$



## Comparison of DP and MC

- **Dynamic Programming**

- **Pros: bootstrapping** (update value estimates based on other learned estimates, without waiting for a final outcome)
- **Cons: Model-Based** (requires complete and accurate model of the environment)

- **Monte Carlo method**

- **Pros: Model-free** (learn directly from experience)
- **Cons:** applicable to **episodic tasks** (require episodes to end to calculate returns, making them unsuitable for continuous tasks)

➔ **Temporal Difference Learning (TD) takes advantages from DP and MC; It is model-free** (uses random sampling), which learns directly from raw experience without a model of the environment, and **bootstraps** by updating estimates based in part on other learned estimates, without waiting for a final outcome.



# Temporal Difference Learning

## Temporal Difference Learning

Value function

$$V_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$$

Using Monte Carlo method

$$V(S_t) \leftarrow V(S_t) + \alpha[G_t - V(S_t)]$$

The episode should be finished to get the return  $G_t$

Total expected return  $G_t$

$$\begin{aligned} G_t &= \sum_{i=t+1}^T \gamma^{i-t-1} R_i \\ &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \\ &= R_{t+1} + \gamma G_{t+1} \end{aligned}$$

State value function

$$\begin{aligned} V_{\pi}(s) &= \mathbb{E}_{\pi}[G_t | S_t = s] \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma G_{t+1} | S_t = s] \\ &= R_{t+1} + \gamma v_{\pi}(S_{t+1}) \end{aligned}$$

## Temporal Difference (TD) Learning

$$V(S_t) \leftarrow V(S_t) + \alpha[\underbrace{R_{t+1} + \gamma V(S_{t+1})}_{\text{TD Target}} - \underbrace{V(S_t)}_{\text{TD Error}}]$$


Estimated value of the current state

- It updates the value of the current state towards the estimated return
- TD error: measures the difference between the predicted value of the current state and the observed reward plus the estimated value of the next state

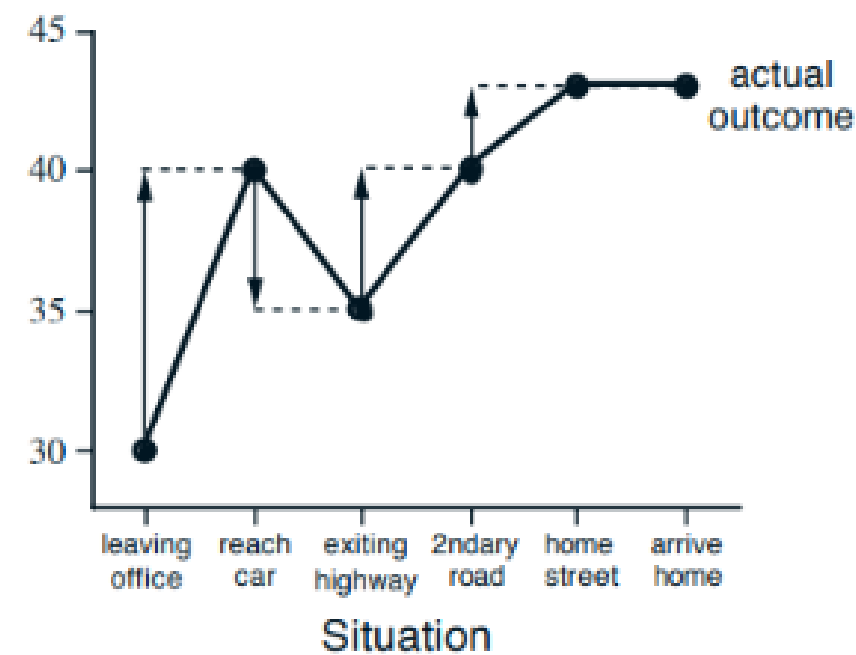
# Temporal Difference Learning

## Temporal Difference Learning

<i>State</i>	<i>Elapsed Time (minutes)</i>	<i>Predicted Time to Go</i>	<i>Predicted Total Time</i>
leaving office, friday at 6	0	30	30
reach car, raining	5	35	40
exiting highway	20	15	35
2ndary road, behind truck	30	10	40
entering home street	40	3	43
arrive home	43	0	43



Monte-Carlo method



Temporal Difference learning

# SARSA

## SARSA

-On-policy TD control

$$S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}$$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \underbrace{(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))}_{\text{TD Error}}$$

Sarsa (on-policy TD control) for estimating  $Q \approx q_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize  $S$

Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

Loop for each step of episode:

Take action  $A$ , observe  $R, S'$

Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)

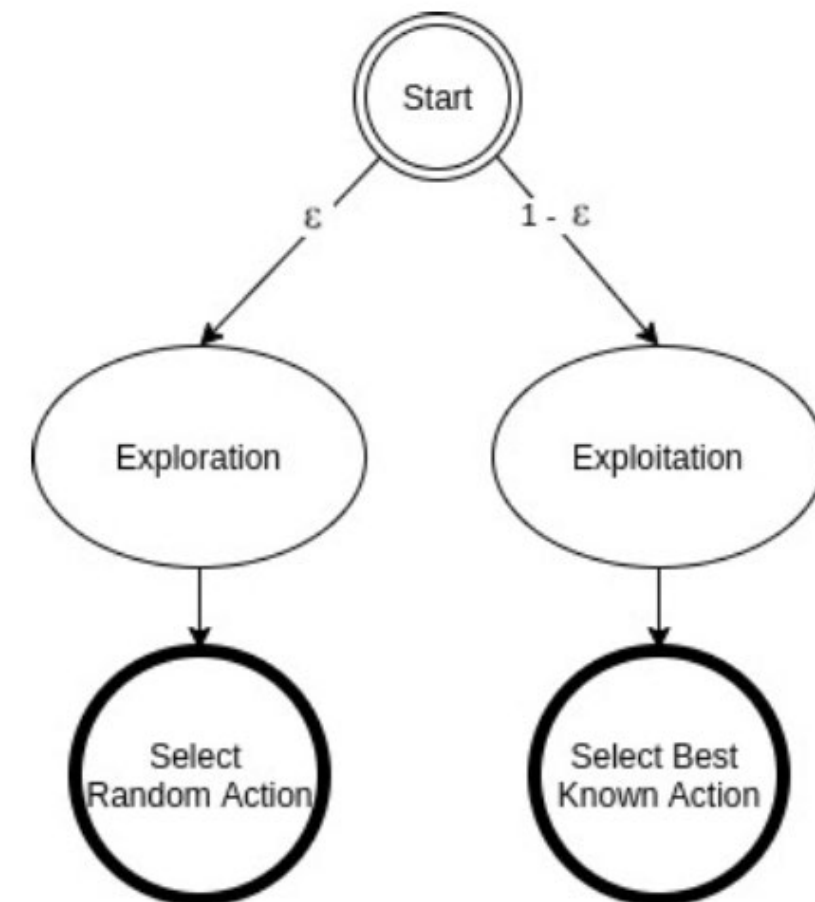
$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until  $S$  is terminal

$\epsilon$ -greedy

$$A \leftarrow \begin{cases} \arg \max_a Q(S, A) & \text{with probability } 1 - \epsilon \\ \text{a random action} & \text{with probability } \epsilon \end{cases}$$



$\epsilon$ -greedy

# Q-learning

## Q-learning

-Off-policy TD control


Update rule in SARSA

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

Update rule in Q-learning

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \underbrace{\max_{a'} Q(S_{t+1}, a')}_{\text{From } \pi_*} - Q(S_t, A_t))$$

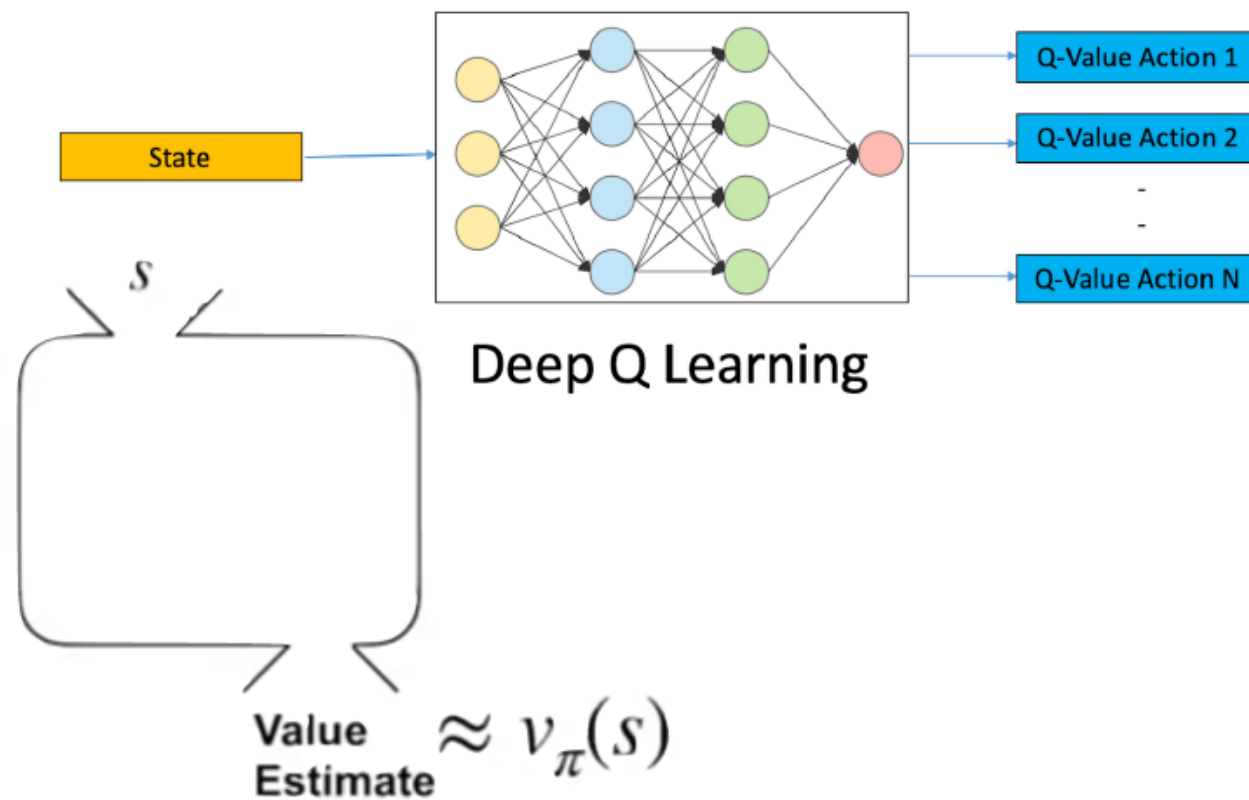
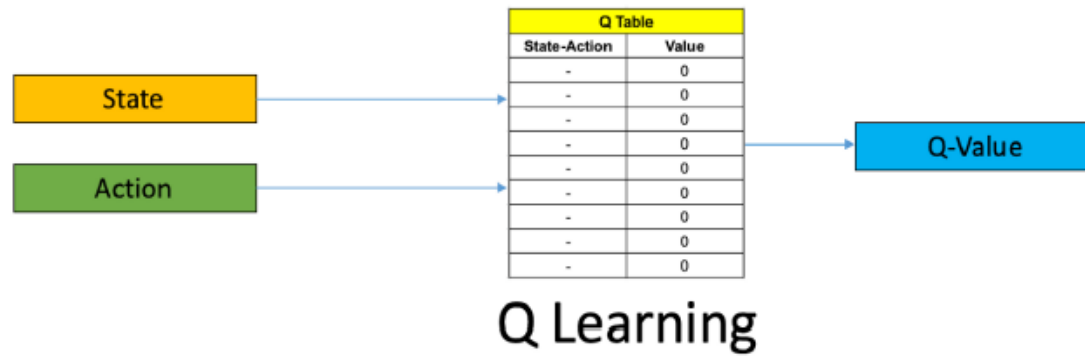
*Target policy  $\pi_*$*



Target and behavior policies

# Q-learning

## Q-learning



Parameterized value function


$$\hat{v}(s, \mathbf{w}) \approx v_{\pi}(s)$$

Linear value function approx.

$$\hat{v}(s, \mathbf{w}) \doteq \sum w_i x_i(s)$$

$$= \langle \mathbf{w}, \mathbf{x}(s) \rangle$$

Nonlinear value function approx. w/ NN






# Deep Q-Networks (DQN)

## Deep Q-Networks (DQN)

### Replay buffer




$s_t, a_t, r_{t+1}, s_{t+1} \rightarrow$


To remove correlations btw samples.

Update with fixed target network  $\theta^-$

$$l = \left( r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta) \right)^2$$



DQN on Atari (Network Architecture)



Pong



Enduro



Beamrider



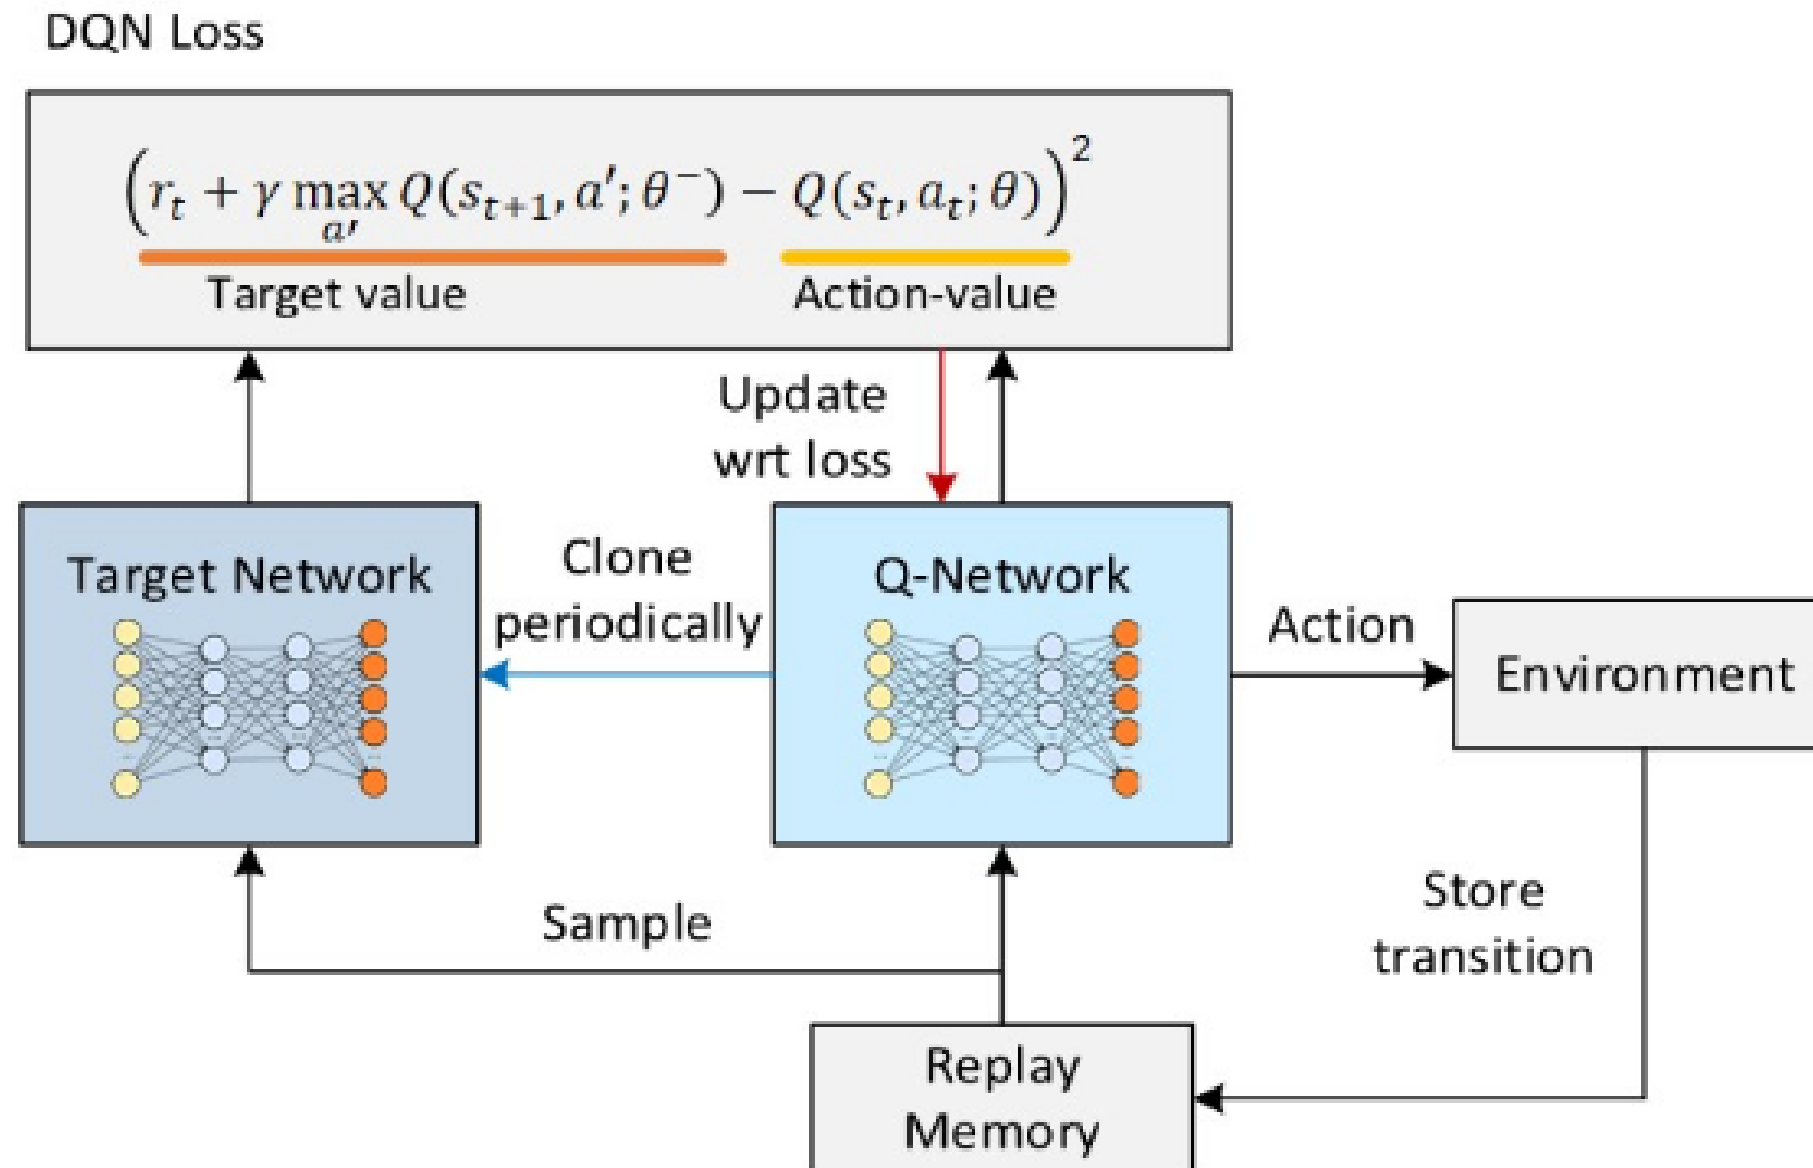
Q\*bert

Atari games

# Deep Q-Networks (DQN)

## Deep Q-Networks (DQN)

- Target network



# Policy Gradient Theorem

## Policy Gradient Theorem

Policy

$$\pi(a | s) = \text{Pr}(A_t = a | S_t = s)$$

↓  
 $\pi_\theta$

Reward function (policy objective function)

$$J(\theta) = \sum_{s \in S} d^\pi(s) V^\pi(s) = \sum_{s \in S} d^\pi(s) \sum_{a \in A} \pi_\theta(a | s) Q^\pi(s, a)$$

Stationary distribution ( $d^\pi$ )

$$d^\pi(s) = \lim_{t \rightarrow \infty} P(s_t = s | s_0, \pi_\theta)$$

*Starting from  $s_0$ , a probability that the state becomes  $s$ .  
(under the policy  $\pi_\theta$ )*

